

Data Structures. Data Structures in C.

Pepe Gallardo, 2024.

Dpto. Lenguajes y Ciencias de la Computación.

University of Málaga

Modules in C

- A non-trivial program is usually divided into several **modules**.
- Each module is a separate **source file** containing a set of related functions.
- Modules are compiled separately and then linked together.
- Modules are usually defined in **header files** (.h) and implemented in **source files** (.c).
- Header files contain the declarations of the public functions, types and global variables provided by the module.
- Source files contain the implementation of the functions, types and variables declared in the header file and any other private declarations needed for the implementation.

Linked Structures

Linked Structures

- **Definition:** A linked structure is a data structure where elements are stored in nodes.
- **Node Composition:** Each node contains a value (element) and a pointer to the next node in the linked structure.
- **End of structure:** The last node has a `NULL` pointer, indicating the end of the structure.
- **Access:** The structure is represented by a [pointer to the first node](#). Starting from this pointer, we can access all nodes by following the pointers.

Advantages:


- **Dynamic Size:** The structure can grow and shrink as needed.
- **Efficient Insertions/Deletions:** No need to move elements around.

Disadvantages:

- **Slower Access:** Accessing arbitrary elements is slower compared to arrays ($O(n)$ vs $O(1)$).
- **Memory Overhead:** Each node requires additional memory for the pointer to the next node.

Example:

- The following picture shows a linked structure of integers:




Definition of a Linked Structure in C

- The following C definition is used to represent a linked structure of integers:
 - `struct Node` represents a node in the linked structure. Each node contains:
 - An integer element.
 - A pointer to the next node in the linked structure.

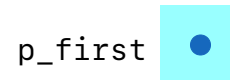
```
struct Node {  
    int element;  
    struct Node* p_next;  
};
```

- Integer `element` in the struct is represented in green/red and `p_next` pointer in blue in the following picture:




Representation of an Empty Linked Structure

- An empty linked structure is represented by a `NULL` pointer, as the first node does not exist:



- A non-empty linked structure is represented by a pointer to its first node. Notice that a `NULL` pointer is also used as a "end of structure" marker in the `p_next` component of the last node:



The `LinkedListStructure.h` Header File

- Notice that functions that modify the linked structure receive a [pointer to the pointer](#) to the first node. This is necessary because the pointer to the first node may change when the structure is modified.

```
#ifndef LINKEDSTRUCTURE_H // Conditional inclusion
#define LINKEDSTRUCTURE_H // Avoids multiple inclusion

#include <stdbool.h>
#include <stddef.h>

// A Node in the linked structure
struct Node {
    int element;
    struct Node* p_next;
};

struct Node* LinkedListStructure_new();
struct Node* LinkedListStructure_copyOf(const struct Node* p_first);
bool LinkedListStructure_isEmpty(const struct Node* p_first);
size_t LinkedListStructure_size(const struct Node* p_first);
void LinkedListStructure_clear(struct Node** p_p_first);
void LinkedListStructure_append(struct Node** p_p_first, int element);
void LinkedListStructure_prepend(struct Node** p_p_first, int element);
void LinkedListStructure_insert(struct Node** p_p_first, size_t index, int element);
int LinkedListStructure_get(const struct Node* p_first, size_t index);
void LinkedListStructure_set(struct Node* p_first, size_t index, int element);
void LinkedListStructure_delete(struct Node** p_p_first, size_t index);
void LinkedListStructure_print(const struct Node* p_first);
#endif
```

The `LinkedListStructure.c` Source File

```
#include <assert.h>
#include <stdio.h>
#include <stdlib.h>

#include "LinkedListStructure.h"

/// @brief Constructor for a new Node
/// @param element Element to be stored in node
/// @param p_next Pointer to next to be stored in node
/// @return Pointer to new node
static struct Node* Node_new(int element, struct Node* p_next) {
    struct Node* p_node = malloc(sizeof(struct Node));
    assert(p_node != NULL && "Node_new: not enough memory");

    p_node->element = element;
    p_node->p_next = p_next;
    return p_node;
}

/// @brief Destructor for a Node. Frees memory allocated for node and sets pointer to node to NULL
/// @param p_p_node Pointer to pointer to node to be freed
static void Node_free(struct Node** p_p_node) {
    free(*p_p_node);
    *p_p_node = NULL;
}

struct Node* LinkedListStructure_new() {
    return NULL;
}

... // rest of the implementation ...
```


Computing the Size of a Linked Structure

- The **size** of a linked structure is the number of nodes it contains.
- The size is computed by traversing the structure and counting the number of nodes.
- We start at the first node and traverse the structure by following the `p_next` pointers, until we reach the end of the structure (a node whose `p_next` pointer is `NULL`).
- We count the number of nodes visited during the traversal.


```
size_t LinkedStructure_size(const struct Node* p_first) {  
    size_t size = 0;  
    ... to complete  
    return size;  
}
```

Getting the Element at a Given Position in a Linked Structure


- We are going to `get` the element at position 2 in the following linked structure:




- Using a pointer (`p_current`), we start at the first node and traverse the structure until we reach the node at the desired position:



- We then return the element stored in that node:




- The `p_current` pointer is an automatic variable so it is deallocated when the function returns:



Setting the Element at a Given Position to a New Value in a Linked Structure


- The same algorithm as for `get` is used to reach the node at the desired position.
- We then update the element stored in that node through the pointer `p_current` :




- For both `get` and `set` operations, it's important to handle **invalid positions** properly.

Prepending an Element at the Beginning of an Empty Linked Structure


- Starting from an empty linked structure, we are going to **prepend** element **5** at the beginning:




- A new node is allocated with element **5** and its **p_next** pointer is set to **NULL**:



- The pointer to the first node is updated to point to the new node:




- p_node** is an automatic variable so it is deallocated when the function returns:




Prepending an Element at the Beginning of a Non-Empty Linked Structure


- Starting from this configuration, we are going to `prepend` element 5 at the beginning of this linked structure:




- A new node is allocated with element 5 and its `p_next` pointer is set to the current first node:



- The pointer to the first node is updated to point to the new node:




- `p_node` is an automatic variable so it is deallocated when the function returns:




Appending an Element at the End of an Empty Linked Structure


- Starting from an empty linked structure, we are going to `append` element **5** at the end:




- A new node is allocated with element **5** and its `p_next` pointer is set to `NULL`:



- The pointer to the first node is updated to point to the new node:



- `p_node` is an automatic variable so it is deallocated when the function returns:




Appending an Element at the End of a Non-Empty Linked Structure


- Starting from this configuration, we are going to **append** element **5** at the end of this linked structure:




- Using a pointer (`p_last`), we start at the first node and traverse the structure until we reach the last node (its `p_next` pointer is `NULL`):




- A new node is allocated with element **5** and its `p_next` pointer is set to `NULL`:




Appending an Element at the End of a Non-Empty Linked Structure (II)



- The `p_next` pointer of the last node is updated to point to the new node:



- `p_last` and `p_node` are automatic variables so they are deallocated when the function returns:



Inserting an Element at an Arbitrary Position in a Linked Structure


- Use pointer `p_current` to traverse the structure to the desired position.
- Use pointer `p_previous` to always track the node before `p_current`.
- Initially, `p_previous` is set to `NULL` and `p_current` points to the first node.
- On each iteration, update `p_previous` to the current node before moving `p_current` to the next node.
- The code to move `p_current` to the desired position (`index`) is as follows:

```
struct Node* p_previous = NULL;
struct Node* p_current = p_first;


for (size_t i = 0; i < index; i++) {
    assert(p_current != NULL && "Invalid index");
    p_previous = p_current;
    p_current = p_current->p_next;
}
```

Inserting an Element at an Arbitrary Position in a Linked Structure (II)


- Starting from this configuration, we are going to insert element 5 at position 2:




- We start at the first node and traverse the structure to move `p_current` to the node at position 2 and `p_previous` to the node before it:




- A new node is allocated with the new element (5) and its `p_next` pointer is set to the node pointed by `p_current`:




Inserting an Element at an Arbitrary Position in a Linked Structure (III)



- The `p_next` pointer of the node pointed by `p_previous` is updated to point to the new node:




- `p_current`, `p_previous` and `p_node` are automatic variables so they are deallocated when the function returns:




Inserting an Element at an Arbitrary Position in a Linked Structure (IV)

Special case: inserting at position 0


- If we want to insert an element at position 0, the loop will not be executed and `p_previous` will be `NULL` and `p_current` will point to the first node:



- As usual, we allocate a new node with the element to insert (5) and set its `p_next` pointer to the node pointed by `p_current`:




- In this case, we also need to update the pointer to the first node to point to the new node:



Inserting an Element at an Arbitrary Position in a Linked Structure (V)

Special case: inserting after the last node


- If we want to insert an element at the end of the structure, the loop will traverse the entire structure, `p_previous` will point to the last node and `p_current` will be `NULL`.
- In this example we are going to insert element **5** at position 3:




- In this case, the same algorithm that we used before in the general case can be used (you should verify that it works correctly).

Deleting an Element at an Arbitrary Position in a Linked Structure


- Starting from this configuration, we are going to delete the element at position 2:




- We start at the first node and traverse the structure until we reach the node at position 2 with `p_current` and the previous node with `p_previous`:




- The `p_next` pointer of the node pointed by `p_previous` is updated to point to the node pointed by the `p_next` component of `p_current`:




Deleting an Element at an Arbitrary Position in a Linked Structure (II)



- The node pointed by `p_current` is deallocated using `free` :




- `p_current` and `p_previous` are automatic variables so they are deallocated when the function returns:




Deleting an Element at an Arbitrary Position in a Linked Structure (III)

Special case: deleting the first node


- If we want to delete the element at position 0, the loop will not be executed and `p_previous` will be `NULL` and `p_current` will point to the first node:



- In this case, we need to update the pointer to the first node to point to the node pointed by the `p_next` component of `p_current`:




- And the node pointed by `p_current` must be deallocated using `free`:



Deleting an Element at an Arbitrary Position in a Linked Structure (IV)

Special case: deleting the last node


- If we want to delete the last node, the loop will traverse the entire structure, `p_previous` will point to the node before the last node and `p_current` will point to the last node.
- In this example we are going to delete the element at position 2:



- In this case, the same algorithm that we used before in the general case can be used (you should verify that it works correctly).

Sorted Linked Structures

- A sorted linked structure is a linked structure where the elements are always stored in some specific order.
- Most of the operations are similar to those of a regular linked structure, except for the insertion.
- The insertion operation takes an element and inserts it in the correct position to **maintain the sorted order**.
- Here is an example of a sorted linked structure of integers in ascending order:



Sorted Linked Structures. Exercise

Create a function named `sorted_insert` that accepts a pointer to a pointer to the first node of a sorted linked structure of integers and an integer element to insert. The function should insert the element in the correct position to maintain the structure's ascending order.

Hint: Utilize the `p_previous` and `p_current` pointers technique to find the appropriate insertion point for the new element.

The `LinkedList` Data Structure

- Some **disadvantages** of representing a linked structure by a pointer to the first node are:
 - We need to pass a pointer to the pointer to the first node to every function that may modify the beginning of the structure. Working with pointers to pointers can be error-prone.
 - Computing the size of the structure requires traversing the entire structure ($O(n)$ complexity).
 - Accessing the last node requires traversing the entire structure ($O(n)$ complexity).

The `LinkedList` Data Structure (II)

- To address these issues, we can define a `LinkedList` data structure that includes a pointer to the first node, a pointer to the last node, and a variable to track the size of the structure.


```
// Node in a LinkedList
struct Node {
    int element;
    struct Node* p_next;
};

// LinkedList implementation
struct LinkedList {
    struct Node* p_first;
    struct Node* p_last;
    size_t size;
};
```

- We will pass a pointer to a `LinkedList` structure to functions that operate on the linked structure. This way, we can access and manipulate the first node, last node, and size of the structure directly.

The LinkedList Data Structure (III)


- An example of a `LinkedList` containing 4 integers is shown below:



- **Invariants:**
 - The `p_first` pointer points to the first node.
 - The `p_last` pointer points to the last node.
 - Each node's `p_next` pointer points to the next node in the structure.
 - The last node's `p_next` pointer is `NULL`.
 - The `size` variable is the number of elements in the structure.

The `LinkedList` Data Structure (IV)

- An example of an empty `LinkedList` is shown below:



The `LinkedList.h` Header File

- All functions now receive a pointer to a `LinkedList`. The only exception is `LinkedList_free`, which receives a pointer to a pointer to a `LinkedList` to allow the function to set the pointer to `NULL`, after freeing the structure, in order to avoid dangling pointers.

```
#ifndef LinkedList_H // conditional compilation directive
#define LinkedList_H // avoids multiple inclusion of the header file

#include <stdbool.h>
#include <stddef.h>

struct Node { // Node in a LinkedList
    int element;
    struct Node* p_next;
};

struct LinkedList { // LinkedList implementation
    struct Node* p_first;
    struct Node* p_last;
    size_t size;
};

struct LinkedList* LinkedList_new();
void LinkedList_free(struct LinkedList** p_p_list);
struct LinkedList* LinkedList_copyOf(const struct LinkedList* p_list);
bool LinkedList_isEmpty(const struct LinkedList* p_list);
size_t LinkedList_size(const struct LinkedList* p_list);
void LinkedList_clear(struct LinkedList* p_list);
void LinkedList_append(struct LinkedList* p_list, int element);
void LinkedList_prepend(struct LinkedList* p_list, int element);
void LinkedList_insert(struct LinkedList* p_list, size_t index, int element);
int LinkedList_get(const struct LinkedList* p_list, size_t index);
void LinkedList_set(const struct LinkedList* p_list, size_t index, int element);
void LinkedList_delete(struct LinkedList* p_list, size_t index);
void LinkedList_print(const struct LinkedList* p_list);
#endif
```


The LinkedList.c File

```
#include <assert.h>
#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>

#include "LinkedList.h"

// Constructor for a new Node
static struct Node* Node_new(int element, struct Node* p_next) {
    struct Node* p_node = malloc(sizeof(struct Node));
    assert(p_node != NULL && "Node_new: not enough memory");

    p_node->element = element;
    p_node->p_next = p_next;
    return p_node;
}

// Destructor for a Node
static void Node_free(struct Node** p_p_node) {
    free(*p_p_node);
    *p_p_node = NULL;
}

// Constructor for a new empty LinkedList
struct LinkedList* LinkedList_new() {
    struct LinkedList* p_list = malloc(sizeof(struct LinkedList));
    assert(p_list != NULL && "LinkedList_new: not enough memory");

    p_list->p_first = NULL;
    p_list->p_last = NULL;
    p_list->size = 0;
    return p_list;
}

... // rest of the implementation
```


The `LinkedList` Data Structure. Exercise

- Complete the implementation of the `LinkedList` data structure by implementing the functions declared in the `LinkedList.h` header file and any other private functions needed.
- Each operation should be implemented in a way that maintains the integrity (**invariants**) of the linked structure: `p_first` should point to the first node in the linked structure, `p_last` should point to the last node, and `size` should be the number of elements in the structure.
- Test the implementation by writing a program that uses the `LinkedList` data structure to store a list of integers and performs various operations on it.

ArrayList

The ArrayList Data Structure

- An `ArrayList` is a dynamic array that can grow to accommodate new elements as needed.
- Initially, it has the ability to hold a certain number of elements, known as its `capacity`. However, this capacity is not fixed and can be expanded dynamically as required.
- Each element's position in the `ArrayList` corresponds directly to its index in the array, with the element at list index `i` being placed at array index `i`.
- Inserting new elements into the `ArrayList` prompts a rightward **shift** of subsequent array elements, creating the necessary space.
- Conversely, when an element is removed, the subsequent elements **shift** leftward, eliminating any gaps.
- The implementation also maintains an integer variable `size` to keep track of the number of elements in the `ArrayList`.
- We are going to define a module for an `ArrayList` of integers. Its implementation looks like this:



The ArrayList.h Header File

```
#ifndef ARRAYLIST_H // conditional compilation directive
#define ARRAYLIST_H // avoids multiple inclusion of the header file

#include <stdbool.h>
#include <stddef.h>

// Structure corresponding to an ArrayList
struct ArrayList {
    int* elements;    // heap allocated array of elements
    size_t capacity;  // capacity of array
    size_t size;      // number of elements in array
};

struct ArrayList* ArrayList_new(size_t initialCapacity);
void ArrayList_free(struct ArrayList* p_list);
struct ArrayList* ArrayList_copyOf(const struct ArrayList* p_list);
bool ArrayList_isEmpty(const struct ArrayList* p_list);
size_t ArrayList_size(const struct ArrayList* p_list);
void ArrayList_clear(struct ArrayList* p_list);
void ArrayList_append(struct ArrayList* p_list, int element);
void ArrayList_prepend(struct ArrayList* p_list, int element);
void ArrayList_insert(struct ArrayList* p_list, size_t index, int element);
int ArrayList_get(const struct ArrayList* p_list, size_t index);
void ArrayList_set(const struct ArrayList* p_list, size_t index, int element);
void ArrayList_delete(struct ArrayList* p_list, size_t index);
void ArrayList_print(const struct ArrayList* p_list);
#endif
```

The ArrayList.c File

```
#include <assert.h>
#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>

#include "ArrayList.h"

// Constructor for a new empty ArrayList with the given initial capacity.
struct ArrayList* ArrayList_new(size_t initialCapacity) {
    assert(initialCapacity > 0 && "ArrayList_new: initialCapacity must be greater than 0");

    struct ArrayList* p_list = malloc(sizeof(struct ArrayList));
    assert(p_list != NULL && "ArrayList_new: not enough memory");


    p_list->elements = malloc(sizeof(int) * initialCapacity);
    assert(p_list->elements != NULL && "ArrayList_new: not enough memory");

    p_list->size = 0;
    p_list->capacity = initialCapacity;
    return p_list;
}

... // rest of the implementation
```


ArrayList After Initialization

- Assuming **initial capacity** is 6, the array of elements has been allocated on the heap. When an array is allocated with `malloc`, its elements will contain garbage values (represented by **?**):




Appending an Element at the End in ArrayList

- `ArrayList_append` adds an element after the last element in the list:
 - Capacity of the array to accommodate the new element must be ensured.
 - The array stores the new element at the index designated by `size`.
 - `size` is incremented to keep track of the number of elements in the list.
- Starting from this configuration, we are going to append element `0` to the `ArrayList` :



- After appending element `0`:



realloc Function in C

- **Purpose:** The `realloc` function is used to resize a memory block that was previously allocated with `malloc` or `calloc`.

- **Declaration:**

```
void* realloc(void* ptr, size_t new_size);
```

- **Functionality:**


- Changes the size of the memory block pointed to by `ptr` to `new_size` bytes.
- Preserves the content of the original memory block up to the minimum of the old and new sizes.
- If `new_size` is larger than the old size, the additional memory is uninitialized.

- **Usage Notes:**


- You do not need to free the original memory block before calling `realloc`.
- The function returns a pointer to the reallocated memory block.
- If reallocation fails, it returns `NULL` and the original memory block remains unchanged.

Ensuring Capacity in ArrayList


- Starting from this configuration, we are going to append element 8 to the ArrayList :



- We reallocate `elements` using `realloc` to double its capacity:

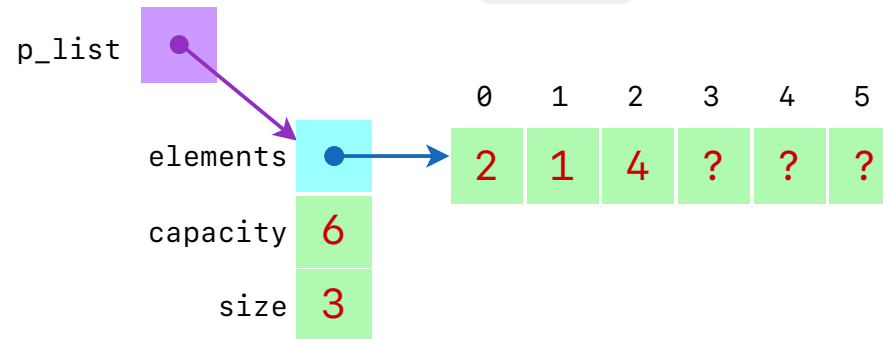


- Now we can append element 8:




Prepending an Element at the Beginning in ArrayList


- Starting from this configuration, we are going to **prepend** element **8** at the beginning:



- The elements at positions 2, 1 and 0 are **shifted rightward** to make space for the new element:




- The new element (**8**) is stored at position 0 and **size** is incremented:




Inserting an Element at an Arbitrary Position in ArrayList


- Starting from this configuration, we are going to **insert** the element **8** at position 1:



- The elements at positions 2 and 1 are **shifted rightward** to make space for the new element:




- The new element (**8**) is stored at position 1 and **size** is incremented:




Deleting an Element at an Arbitrary Position in ArrayList

- Starting from this configuration, we are going to delete the element at position 1:



- The elements at positions 2 and 3 are shifted leftward to fill the gap left by the removed element:



- `size` is decremented :

